



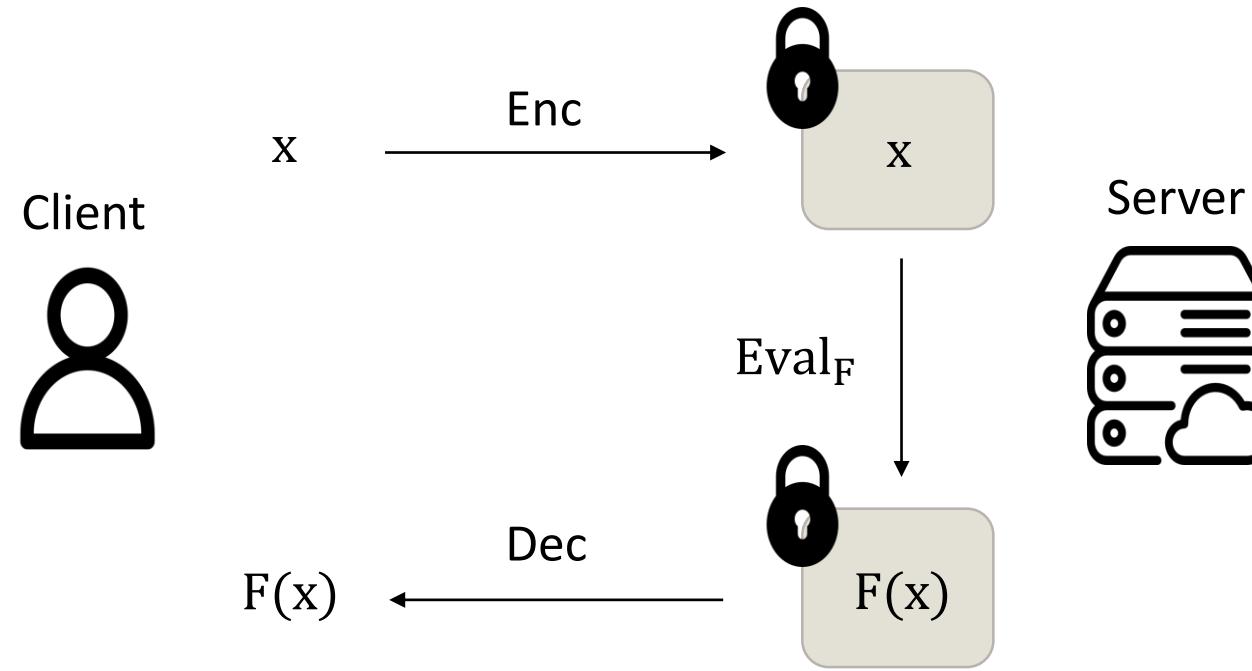
GRAFTING: DECOUPLED SCALE FACTORS AND MODULUS IN RNS-CKKS

Jung Hee Cheon^{1,2}, Hyeongmin Choe², Minsik Kang¹, Jaehyung Kim³

Seonghak Kim², Johannes Mono^{2,4}, Taeyeong Noh²



FULLY HOMOMORPHIC ENCRYPTION (FHE)



- FHE enables computations on encrypted data without decryption.
- Provides efficient privacy-preserving computation.
- CKKS supports approximate computations on real/complex numbers.

RNS-CKKS

- CKKS encodes $\vec{z} \in \mathbb{C}^{N/2}$ with scale factor Δ as:
 - Plaintext: $\Delta m = [\Delta \cdot DFT^{-1}(\vec{z})] \in R_Q = \mathbb{Z}_Q[X]/(X^N + 1)$.

RNS-CKKS

- CKKS encodes $\vec{z} \in \mathbb{C}^{N/2}$ with scale factor Δ as:
 - Plaintext: $\Delta m = [\Delta \cdot DFT^{-1}(\vec{z})] \in R_Q = \mathbb{Z}_Q[X]/(X^N + 1)$.
- CKKS Ciphertext: a pair over $R_Q = \mathbb{Z}_Q[X]/(X^N + 1)$
$$ct(m) = (a, b) \in R_Q^2: \quad a \cdot s + b = \Delta m + e \pmod{Q},$$
where s : secret, e : error, $\Delta m \ll Q$.

RNS-CKKS

- CKKS encodes $\vec{z} \in \mathbb{C}^{N/2}$ with scale factor Δ as:
 - Plaintext: $\Delta m = [\Delta \cdot DFT^{-1}(\vec{z})] \in R_Q = \mathbb{Z}_Q[X]/(X^N + 1)$.
- CKKS Ciphertext: a pair over $R_Q = \mathbb{Z}_Q[X]/(X^N + 1)$
$$ct(m) = (a, b) \in R_Q^2: \quad a \cdot s + b = \Delta m + e \pmod{Q},$$
where s : secret, e : error, $\Delta m \ll Q$.
- For modulus $Q = \prod_{i=0}^{\ell} q_i$, the CRT: $\mathbb{Z}_Q \cong \prod_{i=0}^{\ell} \mathbb{Z}_{q_i}$ allows
$$R_Q \cong R_{q_0} \times R_{q_1} \times \cdots \times R_{q_\ell}$$
 - Computation cost grows linearly with level ℓ .

RNS-CKKS

- CKKS encodes $\vec{z} \in \mathbb{C}^{N/2}$ with scale factor Δ as:
 - Plaintext: $\Delta m = [\Delta \cdot DFT^{-1}(\vec{z})] \in R_Q = \mathbb{Z}_Q[X]/(X^N + 1)$.
 - CKKS Ciphertext: a pair over $R_Q = \mathbb{Z}_Q[X]/(X^N + 1)$
$$ct(m) = (a, b) \in R_Q^2: \quad a \cdot s + b = \Delta m + e \pmod{Q},$$
where s : secret, e : error, $\Delta m \ll Q$.
 - For modulus $Q = \prod_{i=0}^{\ell} q_i$, the CRT: $\mathbb{Z}_Q \cong \prod_{i=0}^{\ell} \mathbb{Z}_{q_i}$ allows
$$R_Q \cong R_{q_0} \times R_{q_1} \times \cdots \times R_{q_\ell}$$
 - Computation cost grows linearly with level ℓ .
- ⇒ Filling Q with machine's word-size primes is the most efficient!

MODULUS RIGIDLY TIED TO SCALE FACTOR

For $ct(m_i) = (a_i, b_i)$ ($i = 1, 2$), CKKS multiplication proceeds as follows:

MODULUS RIGIDLY TIED TO SCALE FACTOR

For $ct(m_i) = (a_i, b_i)$ ($i = 1, 2$), CKKS multiplication proceeds as follows:

- Tensor & Relinearization:

$$\begin{aligned} (d_0, d_1) \in R_Q^2: \quad d_0 + d_1 s &\approx (b_1 + a_1 s)(b_2 + a_2 s) \\ &\approx \Delta^2 m_1 m_2 + \Delta(m_1 e_2 + m_2 e_1) + e_{relin} \pmod{Q}. \end{aligned}$$

MODULUS RIGIDLY TIED TO SCALE FACTOR

For $ct(m_i) = (a_i, b_i)$ ($i = 1, 2$), CKKS multiplication proceeds as follows:

- Tensor & Relinearization:

$$\begin{aligned}(d_0, d_1) \in R_Q^2: \quad d_0 + d_1 s &\approx (b_1 + a_1 s)(b_2 + a_2 s) \\ &\approx \Delta^2 m_1 m_2 + \Delta(m_1 e_2 + m_2 e_1) + e_{relin} \pmod{Q}.\end{aligned}$$

- Rescale:

$$(c_0, c_1) = \left(\left\lfloor \frac{d_0}{q_\ell} \right\rfloor, \left\lfloor \frac{d_1}{q_\ell} \right\rfloor \right) \in R_{Q/q_\ell}^2, \quad c_0 + c_1 s \approx \frac{\Delta^2}{q_\ell} m_1 m_2 + \frac{\Delta}{q_\ell} (m_1 e_2 + m_2 e_1) \pmod{\frac{Q}{q_\ell}}.$$

MODULUS RIGIDLY TIED TO SCALE FACTOR

For $ct(m_i) = (a_i, b_i)$ ($i = 1, 2$), CKKS multiplication proceeds as follows:

- Tensor & Relinearization:

$$\begin{aligned}(d_0, d_1) \in R_Q^2: \quad d_0 + d_1 s &\approx (b_1 + a_1 s)(b_2 + a_2 s) \\ &\approx \Delta^2 m_1 m_2 + \Delta(m_1 e_2 + m_2 e_1) + e_{relin} \pmod{Q}.\end{aligned}$$

- Rescale:

$$(c_0, c_1) = \left(\left\lfloor \frac{d_0}{q_\ell} \right\rfloor, \left\lfloor \frac{d_1}{q_\ell} \right\rfloor \right) \in R_{Q/q_\ell}^2, \quad c_0 + c_1 s \approx \frac{\Delta^2}{q_\ell} m_1 m_2 + \frac{\Delta}{q_\ell} (m_1 e_2 + m_2 e_1) \pmod{\frac{Q}{q_\ell}}.$$

∴ Hence, each modulus should match the scale factor Δ :

$q_1 \approx \dots \approx q_\ell \approx \Delta$ for multiplication levels ℓ .

STRUCTURE ON CKKS MODULUS CHAIN

- Modulus chain in RNS-CKKS is constructed as follows:

$$Q_0 \mid Q_1 \mid \cdots \mid Q_L,$$

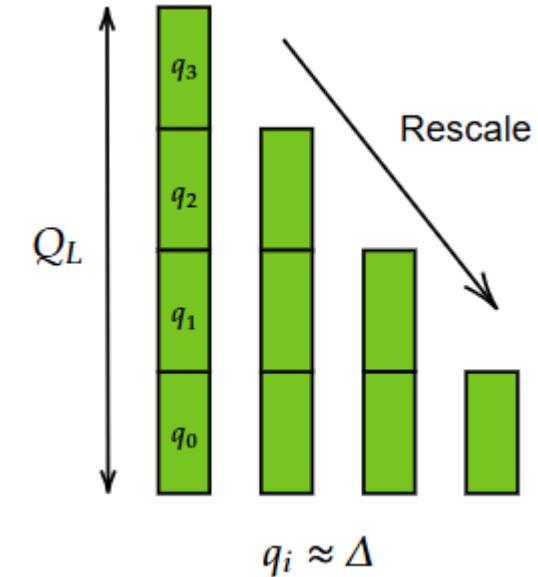
- $Q_\ell = q_0 q_1 \cdots q_\ell \approx \Delta^\ell \cdot q_0$ for each level ℓ .

STRUCTURE ON CKKS MODULUS CHAIN

- Modulus chain in RNS-CKKS is constructed as follows:

$$Q_0 \mid Q_1 \mid \cdots \mid Q_L,$$

- $Q_\ell = q_0 q_1 \cdots q_\ell \approx \Delta^\ell \cdot q_0$ for each level ℓ .



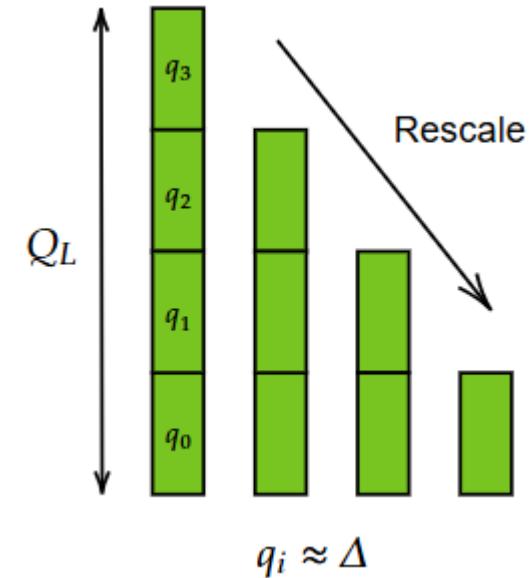
STRUCTURE ON CKKS MODULUS CHAIN

- Modulus chain in RNS-CKKS is constructed as follows:

$$Q_0 \mid Q_1 \mid \cdots \mid Q_L,$$

- $Q_\ell = q_0 q_1 \cdots q_\ell \approx \Delta^\ell \cdot q_0$ for each level ℓ .
- Key-switching keys are positioned at modulus PQ_L
 - The auxiliary modulus P allows key-switching at any level ℓ .
 - The main property enabling this is:

$$Q_\ell \mid Q_L$$



STRUCTURE ON CKKS MODULUS CHAIN

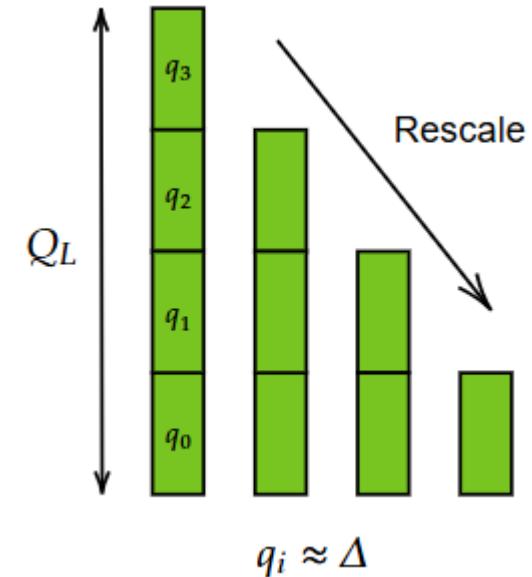
- Modulus chain in RNS-CKKS is constructed as follows:

$$Q_0 \mid Q_1 \mid \cdots \mid Q_L,$$

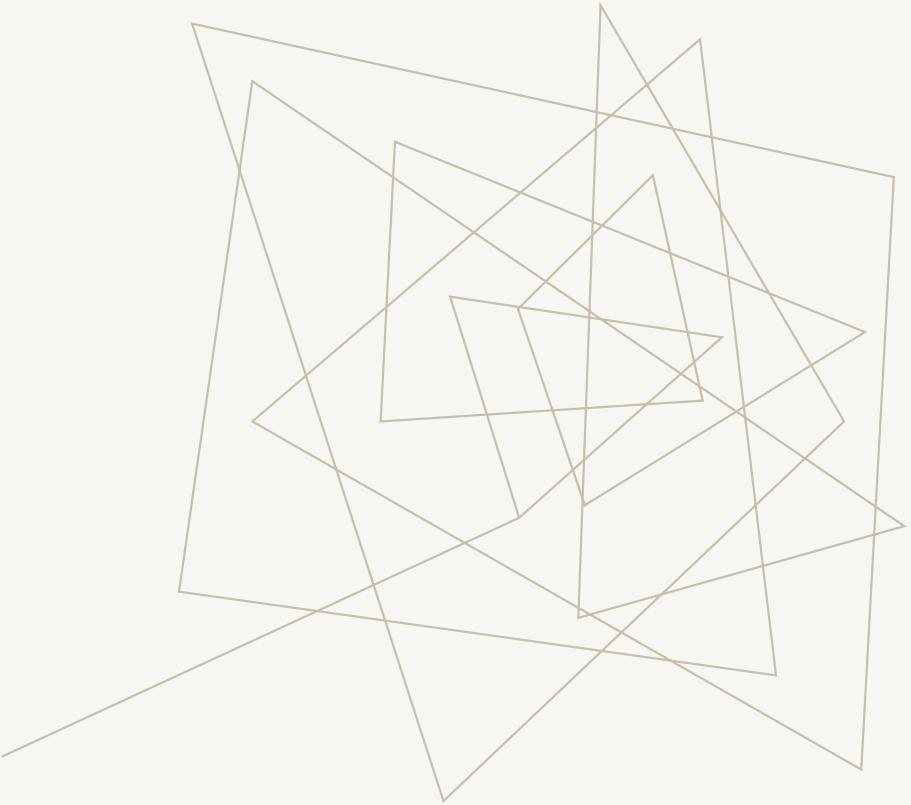
- $Q_\ell = q_0 q_1 \cdots q_\ell \approx \Delta^\ell \cdot q_0$ for each level ℓ .

- Key-switching keys are positioned at modulus PQ_L
 - The auxiliary modulus P allows key-switching at any level ℓ .
 - The main property enabling this is:

$$Q_\ell \mid Q_L$$



Why don't we set $Q_\ell \mid Q_L$ — not necessarily $Q_\ell \mid Q_{\ell+1}$,
while each Q_ℓ is filled up with machine's word-size primes?



GRAFTING: A NOVEL MODULUS MANAGEMENT SYSTEM

RATIONAL RESCALE WITH SPROUT

- We set the top-modulus as $Q_{top} = q_0 q_1 \cdots q_{L-1} \cdot r_{top}$.
 - Each q_i is machine word-size prime.
 - r_{top} is called a *sprout*, reusable modulus factor of Q_{top} .

RATIONAL RESCALE WITH SPROUT

- We set the top-modulus as $Q_{top} = q_0 q_1 \cdots q_{L-1} \cdot r_{top}$.
 - Each q_i is machine word-size prime.
 - r_{top} is called a *sprout*, reusable modulus factor of Q_{top} .
- Every possible modulus is of the form $Q = q_0 q_1 \cdots q_\ell \cdot r$.
 - Here, $\ell < L$ and $r \mid r_{top}$ to ensure $Q \mid Q_{top}$.

RATIONAL RESCALE WITH SPROUT

- We set the top-modulus as $Q_{top} = q_0 q_1 \cdots q_{L-1} \cdot r_{top}$.
 - Each q_i is machine word-size prime.
 - r_{top} is called a *sprout*, reusable modulus factor of Q_{top} .
- Every possible modulus is of the form $Q = q_0 q_1 \cdots q_\ell \cdot r$.
 - Here, $\ell < L$ and $r \mid r_{top}$ to ensure $Q \mid Q_{top}$.
- Rescale from modulus Q to $Q' (< Q)$ proceeds as:

$$\text{Rescale}_{Q \mapsto Q'}(ct = (a, b) \in R_Q^2) = \left(\left\lfloor \frac{Q'}{Q} \cdot a \right\rfloor, \left\lfloor \frac{Q'}{Q} \cdot b \right\rfloor \right) \in R_{Q'}^2,$$

RATIONAL RESCALE WITH SPROUT

- We set the top-modulus as $Q_{top} = q_0 q_1 \cdots q_{L-1} \cdot r_{top}$.
 - Each q_i is machine word-size prime.
 - r_{top} is called a *sprout*, reusable modulus factor of Q_{top} .
- Every possible modulus is of the form $Q = q_0 q_1 \cdots q_\ell \cdot r$.
 - Here, $\ell < L$ and $r \mid r_{top}$ to ensure $Q \mid Q_{top}$.
- Rescale from modulus Q to $Q' (< Q)$ proceeds as:

$$\text{Rescale}_{Q \mapsto Q'}(ct = (a, b) \in R_Q^2) = \left(\left\lfloor \frac{Q'}{Q} \cdot a \right\rfloor, \left\lfloor \frac{Q'}{Q} \cdot b \right\rfloor \right) \in R_{Q'}^2,$$

⇒ We call it *Rational Rescale*, a generalized Rescale in RNS-CKKS.

UNIVERSAL SPROUT

- To enable Rational Rescale with arbitrary bit-lengths,

$$Q = q_0 q_1 \cdots q_\ell \cdot r \text{ should represent integers of any bit-lengths.}$$

UNIVERSAL SPROUT

- To enable Rational Rescale with arbitrary bit-lengths,

$Q = q_0 q_1 \cdots q_\ell \cdot r$ should represent integers of any bit-lengths.

- Sprout r_{top} should satisfy the following:

For word-size ω , $0 \leq \forall \delta < \omega, \exists r \mid r_{top}$ such that $r \approx 2^\delta$.

UNIVERSAL SPROUT

- To enable Rational Rescale with arbitrary bit-lengths,

$Q = q_0 q_1 \cdots q_\ell \cdot r$ should represent integers of any bit-lengths.

- Sprout r_{top} should satisfy the following:

For word-size ω , $0 \leq \forall \delta < \omega, \exists r \mid r_{top}$ such that $r \approx 2^\delta$.

⇒ We call such r_{top} a universal sprout.

UNIVERSAL SPROUT

- To enable Rational Rescale with arbitrary bit-lengths,

$Q = q_0 q_1 \cdots q_\ell \cdot r$ should represent integers of any bit-lengths.

- Sprout r_{top} should satisfy the following:

For word-size ω , $0 \leq \forall \delta < \omega, \exists r \mid r_{top}$ such that $r \approx 2^\delta$.

⇒ We call such r_{top} a universal sprout.

- Example (for $\omega = 60$):

- $r_{top} = 2^{15} \cdot r_1 \cdot r_2$ with r_1 16-bit and r_2 29-bit NTT primes:

UNIVERSAL SPROUT

- To enable Rational Rescale with arbitrary bit-lengths,

$Q = q_0 q_1 \cdots q_\ell \cdot r$ should represent integers of any bit-lengths.

- Sprout r_{top} should satisfy the following:

For word-size ω , $0 \leq \forall \delta < \omega, \exists r \mid r_{top}$ such that $r \approx 2^\delta$.

⇒ We call such r_{top} a universal sprout.

- Example (for $\omega = 60$):

- $r_{top} = 2^{15} \cdot r_1 \cdot r_2$ with r_1 16-bit and r_2 29-bit NTT primes:

	$0 \leq \log_2 r < 16$	$16 \leq \log_2 r < 29$	$29 \leq \log_2 r < 44$	$44 \leq \log_2 r < 60$
r	$2^0, 2^1, \dots, 2^{15}$	$r_1, 2r_1, \dots, 2^{15}r_1$	$r_2, 2r_2, \dots, 2^{15}r_2$	$r_1r_2, 2r_1r_2, \dots, 2^{15}r_1r_2$

UNIVERSAL SPROUT

- To enable Rational Rescale with arbitrary bit-lengths,

$Q = q_0 q_1 \cdots q_\ell \cdot r$ should represent integers of any bit-lengths.

- Sprout r_{top} should satisfy the following:

For word-size ω , $0 \leq \forall \delta < \omega, \exists r \mid r_{top}$ such that $r \approx 2^\delta$.

\Rightarrow We call such r_{top} a universal sprout.

- Example (for $\omega = 60$):

- $r_{top} = 2^{15} \cdot r_1 \cdot r_2$ with r_1 16-bit and r_2 29-bit NTT primes:

	$0 \leq \log_2 r < 16$	$16 \leq \log_2 r < 29$	$29 \leq \log_2 r < 44$	$44 \leq \log_2 r < 60$
r	$2^0, 2^1, \dots, 2^{15}$	$r_1, 2r_1, \dots, 2^{15}r_1$	$r_2, 2r_2, \dots, 2^{15}r_2$	$r_1r_2, 2r_1r_2, \dots, 2^{15}r_1r_2$

1) Embed $\mathbb{Z}_{2^{15}} \hookrightarrow \mathbb{Z}_p$ and 2) composite NTT with $\mathbb{Z}_{r_1} \times \mathbb{Z}_{r_2} \cong \mathbb{Z}_{r_1r_2}$

UNIVERSAL SPROUT

- To enable Rational Rescale with arbitrary bit-lengths,

$Q = q_0 q_1 \cdots q_\ell \cdot r$ should represent integers of any bit-lengths.

- Sprout r_{top} should satisfy the following:

For word-size ω , $0 \leq \forall \delta < \omega, \exists r \mid r_{top}$ such that $r \approx 2^\delta$.

\Rightarrow We call such r_{top} a universal sprout.

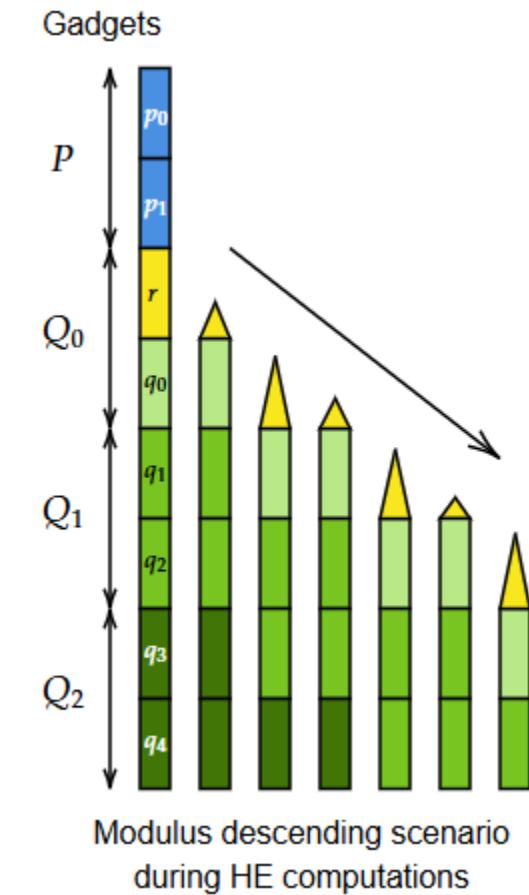
- Example (for $\omega = 60$):

- $r_{top} = 2^{15} \cdot r_1 \cdot r_2$ with r_1 16-bit and r_2 29-bit NTT primes:

	$0 \leq \log_2 r < 16$	$16 \leq \log_2 r < 29$	$29 \leq \log_2 r < 44$	$44 \leq \log_2 r < 60$
r	$2^0, 2^1, \dots, 2^{15}$	$r_1, 2r_1, \dots, 2^{15}r_1$	$r_2, 2r_2, \dots, 2^{15}r_2$	$r_1r_2, 2r_1r_2, \dots, 2^{15}r_1r_2$

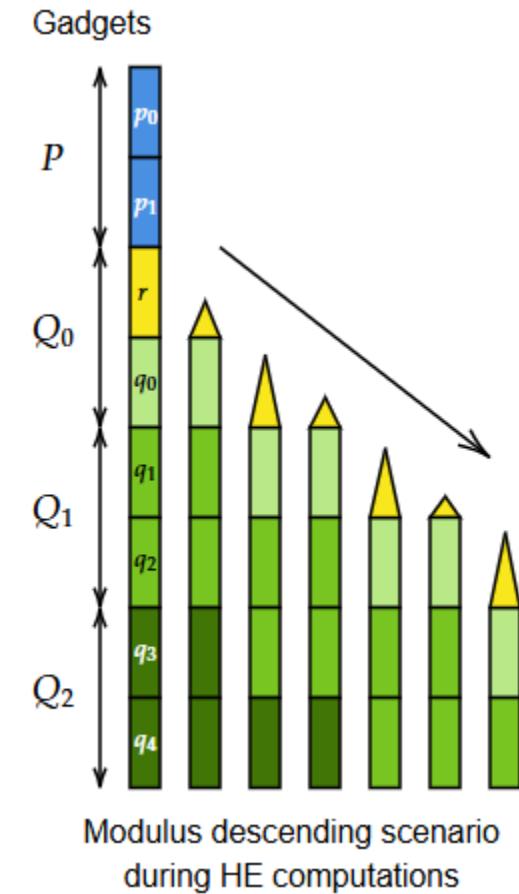
1) Embed $\mathbb{Z}_{2^{15}} \hookrightarrow \mathbb{Z}_p$ and 2) composite NTT with $\mathbb{Z}_{r_1} \times \mathbb{Z}_{r_2} \cong \mathbb{Z}_{r_1r_2}$
 \Rightarrow Universal sprout, within 2 machine words.

MODULUS RESURRECTION DURING COMPUTATIONS



MODULUS RESURRECTION DURING COMPUTATIONS

When rescaling $Q = q_0 q_1 \cdots q_\ell \cdot r$ by q to obtain $Q' \approx Q/q$,



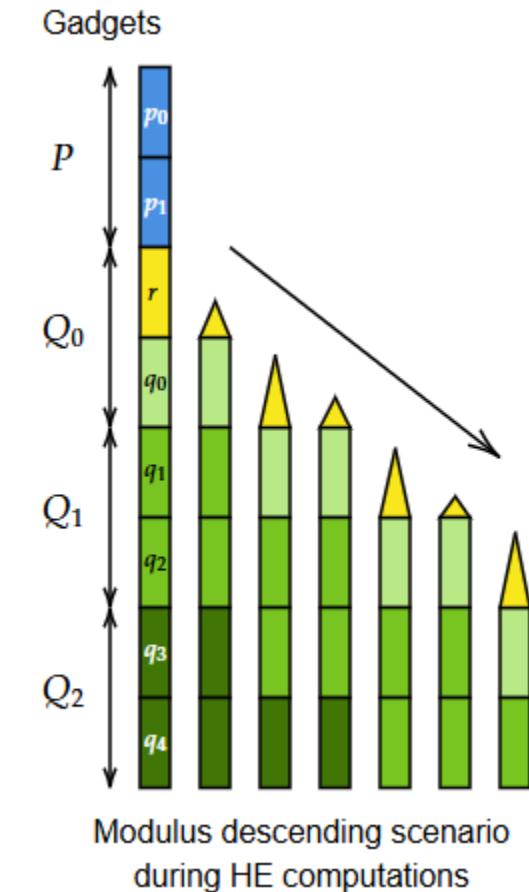
MODULUS RESURRECTION DURING COMPUTATIONS

When rescaling $Q = q_0 q_1 \cdots q_\ell \cdot r$ by q to obtain $Q' \approx Q/q$,

- If $q \leq r$,

$$Q' \approx q_0 q_1 \cdots q_\ell \cdot (r/q).$$

- Choose $r' \mid r_{top}$ such that $r' \approx r/q$.



MODULUS RESURRECTION DURING COMPUTATIONS

When rescaling $Q = q_0q_1 \cdots q_\ell \cdot r$ by q to obtain $Q' \approx Q/q$,

- If $q \leq r$,

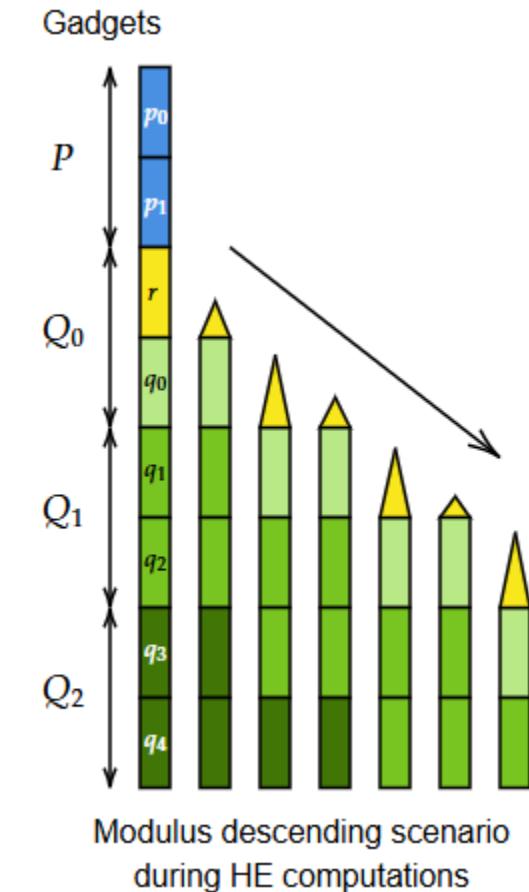
$$Q' \approx q_0q_1 \cdots q_\ell \cdot (r/q).$$

- Choose $r' \mid r_{top}$ such that $r' \approx r/q$.

- If $q > r$,

$$Q' \approx q_0q_1 \cdots q_{\ell-1} \cdot (q_\ell r/q).$$

- Choose $r' \mid r_{top}$ such that $r' \approx (2^\omega r)/q$.
- Some factors of r_{top} are resurrected during Rational Rescale.



MODULUS RESURRECTION DURING COMPUTATIONS

When rescaling $Q = q_0 q_1 \cdots q_\ell \cdot r$ by q to obtain $Q' \approx Q/q$,

- If $q \leq r$,

$$Q' \approx q_0 q_1 \cdots q_\ell \cdot (r/q).$$

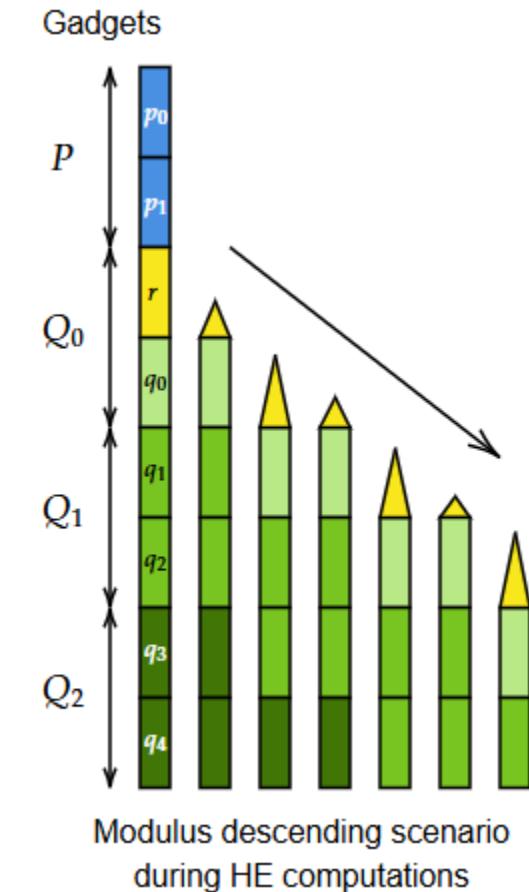
- Choose $r' \mid r_{top}$ such that $r' \approx r/q$.

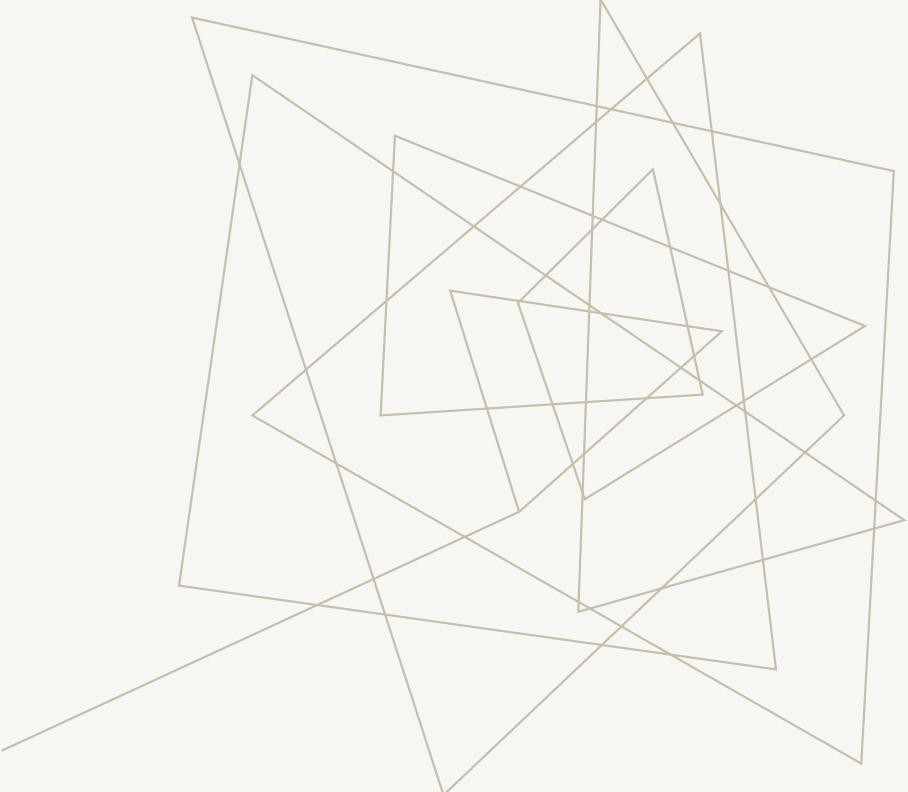
- If $q > r$,

$$Q' \approx q_0 q_1 \cdots q_{\ell-1} \cdot (q_\ell r/q).$$

- Choose $r' \mid r_{top}$ such that $r' \approx (2^\omega r)/q$.
- Some factors of r_{top} are resurrected during Rational Rescale.

⇒ We call this process Modulus Resurrection.





APPLICATION OF GRAFTING & EXPERIMENTAL RESULTS

WHEN APPLIED TO STANDARD CKKS

Parameter	N	log PQ	# factors	Mult (ms)	Boot (s)	Key size (MB)
HEaaN [Cry22]	2^{15}	777	22 ↘	102.20	14.5	115.34
		780	13 ↘	57.28 (1.8x)	7.6 (1.9x)	44.04 (62% ↓)

- CKKS parameters can benefit from fewer RNS factors by Grafting.
- Operations gain up to **2.0x speed-up**, with **key size reduced by up to 62%**.

WHEN APPLIED TO STANDARD CKKS

Parameter	N	log PQ	# factors	Mult (ms)	Boot (s)	Key size (MB)
HEaaN [Cry22]	2^{15}	777	22 ↘	102.20	14.5	115.34
		780	13 ↘	57.28 (1.8x)	7.6 (1.9x)	44.04 (62% ↓)
Sec. Guide. [BCC+24]	2^{16}	1734	35 ↘	360.84	86.5	220.20
			29 ↘	179.87 (2.0x)	71.7 (1.2x)	157.29 (29% ↓)

- CKKS parameters can benefit from fewer RNS factors by Grafting.
- Operations gain up to **2.0x speed-up**, with **key size reduced by up to 62%**.

WHEN APPLIED TO STANDARD CKKS

Parameter	N	log PQ	# factors	Mult (ms)	Boot (s)	Key size (MB)
HEaaN [Cry22]	2^{15}	777	22 ↘	102.20	14.5	115.34
		780	13 ↘	57.28 (1.8x)	7.6 (1.9x)	44.04 (62% ↓)
Sec. Guide. [BCC+24]	2^{16}	1734	35 ↘	360.84	86.5	220.20
			29 ↘	179.87 (2.0x)	71.7 (1.2x)	157.29 (29% ↓)
HEaaN [Cry22]	2^{16}	1555	30 ↘	329.38	37.0	157.29
			27 ↘	247.45 (1.3x)	35.5 (1.1x)	146.80 (7% ↓)

- CKKS parameters can benefit from fewer RNS factors by Grafting.
- Operations gain up to **2.0x speed-up**, with **key size reduced by up to 62%**.

WHEN APPLIED TO STANDARD CKKS

Parameter	N	log PQ	# factors	Mult (ms)	Boot (s)	Key size (MB)
HEaaN [Cry22]	2^{15}	777	22 ↘	102.20	14.5	115.34
		780	13 ↘	57.28 (1.8x)	7.6 (1.9x)	44.04 (62% ↓)
Sec. Guide. [BCC+24]	2^{16}	1734	35 ↘	360.84	86.5	220.20
			29 ↘	179.87 (2.0x)	71.7 (1.2x)	157.29 (29% ↓)
HEaaN [Cry22]	2^{16}	1555	30 ↘	329.38	37.0	157.29
			27 ↘	247.45 (1.3x)	35.5 (1.1x)	146.80 (7% ↓)

- CKKS parameters can benefit from fewer RNS factors by Grafting.
- Operations gain up to **2.0x speed-up**, with **key size reduced by up to 62%**.
 - Parameters with small scale factors → accelerates well!
 - Parameters with large scale factors → not by much

WHEN APPLIED TO BIT-CKKS [BCKS24]

Parameter	N	log PQ	# factors	Mult (ms)	GateBoot (s)	Key size (MB)
Bit-CKKS [BCKS24]	2^{14}	424	14	16.1	5.18	47.71
		426	8	16.2 (1.8x)	2.74 (1.9x)	16.52 (65% ↓)
	2^{16}	1598	46	884.8	102.10	144.70
		1522	25	428.1 (2.1x)	56.41 (1.8x)	109.05 (25% ↓)

- Low-precision parameters or Bit/Discrete-CKKS can also benefit from fewer RNS factors!
 - Bunch of small scale factors → accelerates well!

WHEN APPLIED TO BIT-CKKS [BCKS24]

Parameter	N	log PQ	# factors	Mult (ms)	GateBoot (s)	Key size (MB)
Bit-CKKS [BCKS24]	2^{14}	424	14	16.1	5.18	47.71
		426	8	16.2 (1.8x)	2.74 (1.9x)	16.52 (65% ↓)
	2^{16}	1598	46	884.8	102.10	144.70
		1522	25	428.1 (2.1x)	56.41 (1.8x)	109.05 (25% ↓)

- Low-precision parameters or Bit/Discrete-CKKS can also benefit from fewer RNS factors!
 - Bunch of small scale factors → accelerates well!
- Operations gain up to **2.1x speed-up**, with **key size reduced by up to 65%**.

APPLICATION TO HOMOMORPHIC COMPARISON [CKK20]

- **Δ -Changeability:** Decoupled Δ and Q allows changing Δ freely during computation!

APPLICATION TO HOMOMORPHIC COMPARISON [CKK20]

- **Δ -Changeability:** Decoupled Δ and Q allows changing Δ freely during computation!
- **Application?**
 - Adaptive Precision Computation in **Plain World** allows better latency/memory:
 - ML training: $\text{FP16} \Leftrightarrow \text{FP32}$
 - Iterative solvers: $\text{FP8} \rightarrow \text{FP16} \rightarrow \text{FP32}$

APPLICATION TO HOMOMORPHIC COMPARISON [CKK20]

- **Δ -Changeability:** Decoupled Δ and Q allows changing Δ freely during computation!
- **Application?**
 - Adaptive Precision Computation in **Plain World** allows better latency/memory:
 - ML training: FP16 \Leftrightarrow FP32
 - Iterative solvers: FP8 \rightarrow FP16 \rightarrow FP32
 - In **Encrypted World**, Grafting allows changing **precision** by changing Δ
 - E.g., Iterative solver with $\Delta_0 \ll \Delta_1 \ll \Delta_2$
 - For each $x_{i+1} = f(x_i)$, we can use fine-tuned Δ_i
 - Mod consumption: $3\Delta_2 \cdot depth_f \rightarrow (\Delta_0 + \Delta_1 + \Delta_2) \cdot depth_f$

$$(depth_f = \lceil \log_2 \deg f + 1 \rceil)$$

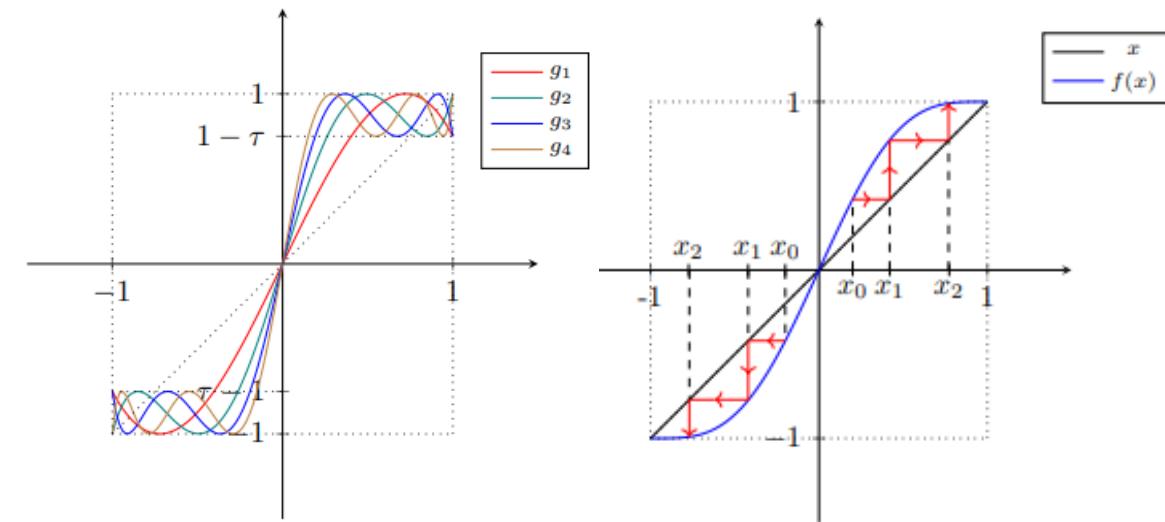
13

APPLICATION TO HOMOMORPHIC COMPARISON [CKK20]

- Homomorphic Comparison [CKK20] use iterative method:

- Evaluate $f^k \circ g^\ell: I_\epsilon \rightarrow I_{1-2^{-\alpha}}$,
 - for low-degree poly f (and g),
 - Iteratively narrowing the interval.

$$(I_\epsilon := [-1, -\epsilon] \cup [\epsilon, 1])$$

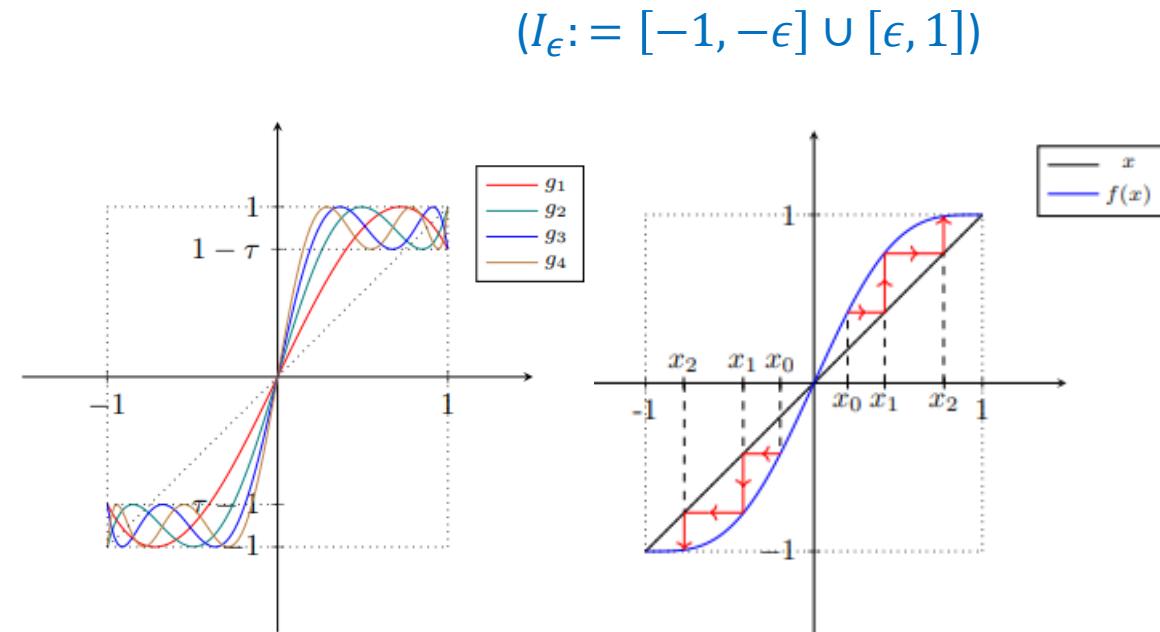
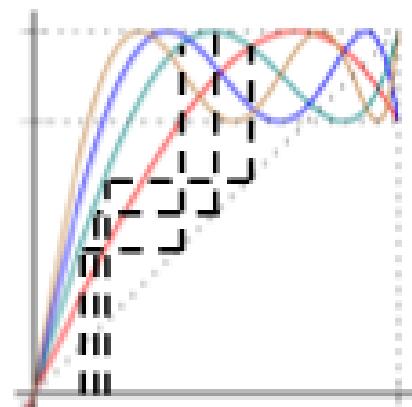


APPLICATION TO HOMOMORPHIC COMPARISON [CKK20]

- Homomorphic Comparison [CKK20] use iterative method:

- Evaluate $f^k \circ g^\ell: I_\epsilon \rightarrow I_{1-2^{-\alpha}}$,
 - for low-degree poly f (and g),
 - Iteratively narrowing the interval.

- First, g narrows the interval, but precision is not very important.

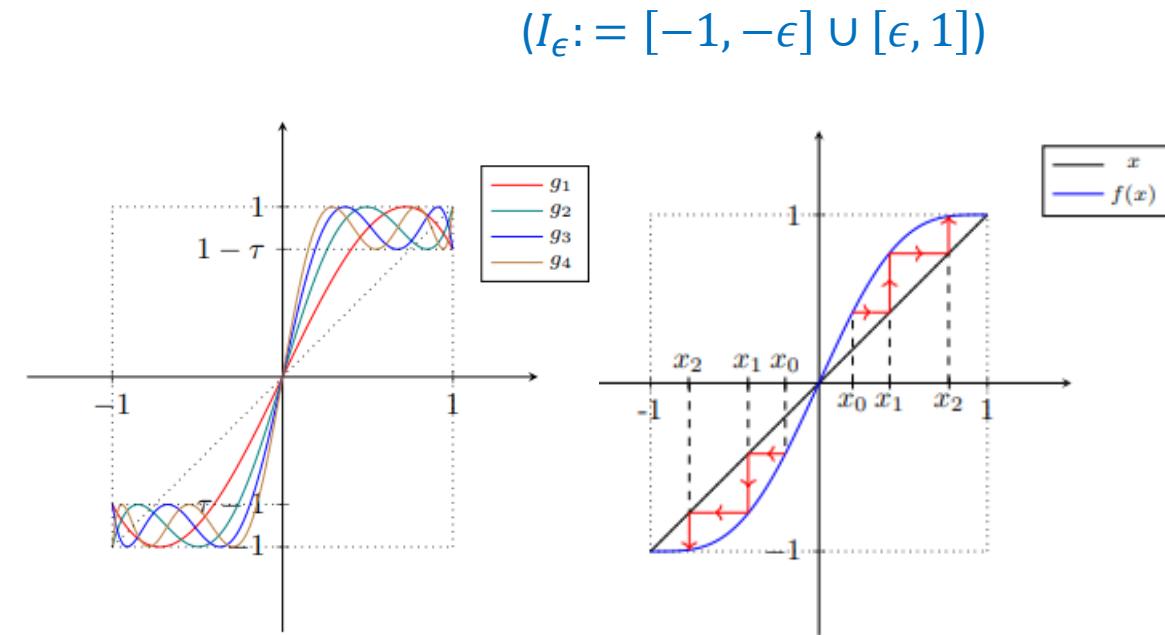
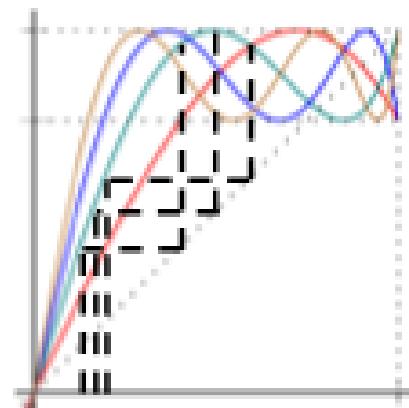


APPLICATION TO HOMOMORPHIC COMPARISON [CKK20]

- Homomorphic Comparison [CKK20] use iterative method:

- Evaluate $f^k \circ g^\ell: I_\epsilon \rightarrow I_{1-2^{-\alpha}}$,
 - for low-degree poly f (and g),
 - Iteratively narrowing the interval.

- First, g narrows the interval, but precision is not very important.



→ We can save modulus by using smaller Δ

APPLICATION TO HOMOMORPHIC COMPARISON [CKK20]

- Homomorphic Comparison [CKK20] use iterative method:
 - We use degree-7 polynomials f_3 and g_3
→ Use depth 3 per polynomial evaluation

APPLICATION TO HOMOMORPHIC COMPARISON [CKK20]

- Homomorphic Comparison [CKK20] use iterative method:
 - We use degree-7 polynomials f_3 and g_3
→ Use depth 3 per polynomial evaluation
 - We can reduce the early scale factors while keeping accuracy!

APPLICATION TO HOMOMORPHIC COMPARISON [CKK20]

- Homomorphic Comparison [CKK20] use iterative method:
 - We use degree-7 polynomials f_3 and g_3
→ Use depth 3 per polynomial evaluation
 - We can reduce the early scale factors while keeping accuracy!

Comparison Function	$f_3^{(2)} \circ g_3^{(4)}$		$f_3^{(2)} \circ g_3^{(8)}$	
Methods	Original	Changing Δ (28, 30, 42 bits)	Original	Changing Δ (31, 42 bits)
Consumed Modulus (bit)	$(42 \times 3) \times 6 = 756$	$(28 \times 3) \times 4 + (30 \times 3) + (42 \times 3) = 552$ (27% ↓)	$(42 \times 3) \times 10 = 1260$	$(31 \times 3) \times 9 + (42 \times 3) = 963$ (24% ↓)
Precision (bit)	23.1	23.1	23.5	23.3

⌘ Bit-precision := $-\log_2 |\max \text{ error}|$ from 100 iterations

SUMMARY

- Grafting redesigns modulus usage in RNS-CKKS
 - Enable arbitrary bit-lengths Rescale,
 - No additional key-switching keys required.

SUMMARY

- Grafting redesigns modulus usage in RNS-CKKS
 - Enable arbitrary bit-lengths Rescale,
 - No additional key-switching keys required.

⇒ Grafting improves RNS-CKKS:

1. Machine word-size RNS Modulus
 - Up to 2.01x **faster** multiplication and bootstrapping.
 - Up to 62% **reduction in** ciphertext/key-switching key **size**.

SUMMARY

- Grafting redesigns modulus usage in RNS-CKKS
 - Enable arbitrary bit-lengths Rescale,
 - No additional key-switching keys required.

⇒ Grafting improves RNS-CKKS:

1. Machine word-size RNS Modulus
 - Up to 2.01x **faster** multiplication and bootstrapping.
 - Up to 62% **reduction in** ciphertext/key-switching key **size**.
2. Flexible Scale factor
 - Scale adjustable independently of the modulus.
 - **Modulus saving** for iterative methods such as in [CKK20].



EPRINT: 2024/1014

THANK YOU!